

# Neural Networks image recognition - MultiLayer Perceptron

Use both MLNN for the following problem.

1. Add random noise (see below on `size` parameter on `np.random.normal`) to the images in training and testing. **Make sure each image gets a different noise feature added to it. Inspect by printing out several images. Note - the `size` parameter should match the data.**
2. Compare the `accuracy` of train and val after N epochs for MLNN with and without noise.
3. Vary the amount of noise by changing the `scale` parameter in `np.random.normal` by a factor. Use `.1, .5, 1.0, 2.0, 4.0` for the `scale` and keep track of the `accuracy` for training and validation and plot these results.

```
np.random.normal
```

## Parameters

### loc

Mean ("centre") of the distribution.

### scale

Standard deviation (spread or "width") of the distribution. Must be non-negative.

### size

Output shape. If the given shape is, e.g., (m, n, k), then m *n* k samples are drawn. If size is None (default), a single value is returned if loc and scale are both scalars. Otherwise, np.broadcast(loc, scale).size samples are drawn.

## Neural Networks - Image Recognition

```
In [102... import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.optimizers import RMSprop
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras import backend

import numpy as np
```

```
In [103... import matplotlib.pyplot as plt
%matplotlib inline
from IPython.display import display
```

## Multi Layer Neural Network

Trains a simple deep NN on the MNIST dataset. Gets to 98.40% test accuracy after 20 epochs (there is *a lot* of margin for parameter tuning).


```
In [104... # the data, shuffled and split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()
```

```
In [105... y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

model = Sequential()
```


```
In [106... print("original, no noise")
plt.imshow(x_train[1])
```

```
original, no noise
Out[106]: <matplotlib.image.AxesImage at 0x2467a19d850>
```




```
In [107]: noisy = x_train + np.random.normal(0,.5,(60000,28,28))
          print("a little noise added")
          plt.imshow(noisy[1])
```

```
a little noise added
Out[107]: <matplotlib.image.AxesImage at 0x24670867b80>
```



```
In [108]: noisy2 = x_train + np.random.normal(0,4,(60000,28,28))
          print("very noisy")
          plt.imshow(noisy2[1])
```

```
very noisy
Out[108]: <matplotlib.image.AxesImage at 0x24670a1e4f0>
```



```
In [109...
scale = [0, 0.1, 0.5, 1.0, 2.0, 4.0]
batch_size = 128
num_classes = 10
epochs = 20
# Noise is added here
# The max value of the noise should not grossly surpass 1.0
for s in scale:
    noise_train = np.random.normal(0, s, size=(60000, 28, 28))
    new_train = noise_train + x_train
    noise_test = np.random.normal(0, s, size=(10000, 28, 28))
    new_test = noise_test + x_test

    new_train = new_train.reshape(60000, 784)
    new_test = new_test.reshape(10000, 784)
    new_train = new_train.astype('float32')
    new_test = new_test.astype('float32')
    new_train /= 255
    new_test /= 255

    model.add(Dense(512, activation='relu', input_shape=(784,)))
    model.add(Dropout(0.2))
    model.add(Dense(512, activation='relu'))
    model.add(Dropout(0.2))
    model.add(Dense(10, activation='softmax'))

    model.summary()
    model.compile(loss='categorical_crossentropy',
                  optimizer="adam",
                  metrics=['accuracy'])

    model.fit(new_train, y_train,
              batch_size=batch_size,
              epochs=epochs,
              verbose=1,
              validation_data=(new_test, y_test))

    score = model.evaluate(new_test, y_test, verbose=0)
    print('Scores for noise scale ' + str(s))
    print('Test loss:', score[0])
    print('Test accuracy:', score[1])
```

Model: "sequential\_5"

Layer (type)	Output Shape	Param #
dense_39 (Dense)	(None, 512)	401920
dropout_26 (Dropout)	(None, 512)	0
dense_40 (Dense)	(None, 512)	262656
dropout_27 (Dropout)	(None, 512)	0
dense_41 (Dense)	(None, 10)	5130

=====  
Total params: 669,706  
Trainable params: 669,706  
Non-trainable params: 0

Epoch 1/20  
469/469 [=====] - 10s 14ms/step - loss: 0.2508 - accuracy: 0.9258 - val\_loss: 0.1025 - val\_accuracy: 0.9685  
Epoch 2/20  
469/469 [=====] - 6s 13ms/step - loss: 0.1020 - accuracy: 0.9685 - val\_loss: 0.0875 - val\_accuracy: 0.9723  
Epoch 3/20  
469/469 [=====] - 5s 10ms/step - loss: 0.0713 - accuracy: 0.9773 - val\_loss: 0.0676 - val\_accuracy: 0.9790  
Epoch 4/20  
469/469 [=====] - 5s 10ms/step - loss: 0.0558 - accuracy: 0.9826 - val\_loss: 0.0742 - val\_accuracy: 0.9771  
Epoch 5/20  
469/469 [=====] - 5s 10ms/step - loss: 0.0443 - accuracy: 0.9854 - val\_loss: 0.0834 - val\_accuracy: 0.9745  
Epoch 6/20  
469/469 [=====] - 5s 10ms/step - loss: 0.0397 - accuracy: 0.9870 - val\_loss: 0.0684 - val\_accuracy: 0.9795  
Epoch 7/20  
469/469 [=====] - 5s 10ms/step - loss: 0.0334 - accuracy: 0.9886 - val\_loss: 0.0713 - val\_accuracy: 0.9805  
Epoch 8/20  
469/469 [=====] - 5s 10ms/step - loss: 0.0300 - accuracy: 0.9899 - val\_loss: 0.0730 - val\_accuracy: 0.9795  
Epoch 9/20  
469/469 [=====] - 5s 11ms/step - loss: 0.0280 - accuracy: 0.9901 - val\_loss: 0.0657 - val\_accuracy: 0.9822  
Epoch 10/20  
469/469 [=====] - 5s 10ms/step - loss: 0.0258 - accuracy: 0.9913 - val\_loss: 0.0761 - val\_accuracy: 0.9796  
Epoch 11/20  
469/469 [=====] - 5s 11ms/step - loss: 0.0216 - accuracy: 0.9924 - val\_loss: 0.0771 - val\_accuracy: 0.9812  
Epoch 12/20  
469/469 [=====] - 5s 10ms/step - loss: 0.0218 - accuracy: 0.9930 - val\_loss: 0.0719 - val\_accuracy: 0.9815  
Epoch 13/20  
469/469 [=====] - 5s 10ms/step - loss: 0.0212 - accuracy: 0.9928 - val\_loss: 0.0764 - val\_accuracy: 0.9821  
Epoch 14/20  
469/469 [=====] - 5s 10ms/step - loss: 0.0190 - accuracy: 0.9936 - val\_loss: 0.0951 - val\_accuracy: 0.9769  
Epoch 15/20  
469/469 [=====] - 5s 10ms/step - loss: 0.0178 - accuracy: 0.9941 - val\_loss: 0.0738 - val\_accuracy: 0.9825  
Epoch 16/20  
469/469 [=====] - 5s 10ms/step - loss: 0.0152 - accuracy: 0.9948 - val\_loss: 0.0811 - val\_accuracy: 0.9826  
Epoch 17/20  
469/469 [=====] - 5s 10ms/step - loss: 0.0203 - accuracy: 0.9932 - val\_loss: 0.0764 - val\_accuracy: 0.9828  
Epoch 18/20  
469/469 [=====] - 5s 10ms/step - loss: 0.0141 - accuracy: 0.9956 - val\_loss: 0.0860 - val\_accuracy: 0.9799  
Epoch 19/20  
469/469 [=====] - 5s 10ms/step - loss: 0.0183 - accuracy: 0.9939 - val\_loss: 0.0801 - val\_accuracy: 0.9832  
Epoch 20/20  
469/469 [=====] - 5s 10ms/step - loss: 0.0143 - accuracy: 0.9954 - val\_loss: 0.0777 - val\_accuracy: 0.9839  
Scores for noise scale 0  
Test loss: 0.07769788056612015  
Test accuracy: 0.9839000105857849  
Model: "sequential\_5"

Layer (type)	Output Shape	Param #
=====	=====	=====

dense_39 (Dense)	(None, 512)	401920
dropout_26 (Dropout)	(None, 512)	0
dense_40 (Dense)	(None, 512)	262656
dropout_27 (Dropout)	(None, 512)	0
dense_41 (Dense)	(None, 10)	5130
dense_42 (Dense)	(None, 512)	5632
dropout_28 (Dropout)	(None, 512)	0
dense_43 (Dense)	(None, 512)	262656
dropout_29 (Dropout)	(None, 512)	0
dense_44 (Dense)	(None, 10)	5130

=====  
Total params: 943,124  
Trainable params: 943,124  
Non-trainable params: 0

---

Epoch 1/20  
469/469 [=====] - 12s 19ms/step - loss: 0.1066 - accuracy: 0.9893 - val\_loss: 0.1364 - val\_accuracy: 0.9789  
Epoch 2/20  
469/469 [=====] - 8s 18ms/step - loss: 0.0554 - accuracy: 0.9906 - val\_loss: 0.1265 - val\_accuracy: 0.9803  
Epoch 3/20  
469/469 [=====] - 6s 14ms/step - loss: 0.0535 - accuracy: 0.9908 - val\_loss: 0.1233 - val\_accuracy: 0.9795  
Epoch 4/20  
469/469 [=====] - 7s 14ms/step - loss: 0.0433 - accuracy: 0.9924 - val\_loss: 0.1075 - val\_accuracy: 0.9820  
Epoch 5/20  
469/469 [=====] - 7s 14ms/step - loss: 0.0367 - accuracy: 0.9932 - val\_loss: 0.0984 - val\_accuracy: 0.9827  
Epoch 6/20  
469/469 [=====] - 7s 14ms/step - loss: 0.0348 - accuracy: 0.9934 - val\_loss: 0.1260 - val\_accuracy: 0.9791  
Epoch 7/20  
469/469 [=====] - 6s 13ms/step - loss: 0.0322 - accuracy: 0.9936 - val\_loss: 0.1269 - val\_accuracy: 0.9804  
Epoch 8/20  
469/469 [=====] - 7s 14ms/step - loss: 0.0300 - accuracy: 0.9938 - val\_loss: 0.0980 - val\_accuracy: 0.9840  
Epoch 9/20  
469/469 [=====] - 7s 15ms/step - loss: 0.0232 - accuracy: 0.9949 - val\_loss: 0.0975 - val\_accuracy: 0.9836  
Epoch 10/20  
469/469 [=====] - 7s 14ms/step - loss: 0.0271 - accuracy: 0.9940 - val\_loss: 0.0980 - val\_accuracy: 0.9832  
Epoch 11/20  
469/469 [=====] - 7s 15ms/step - loss: 0.0292 - accuracy: 0.9935 - val\_loss: 0.0998 - val\_accuracy: 0.9825  
Epoch 12/20  
469/469 [=====] - 7s 14ms/step - loss: 0.0299 - accuracy: 0.9934 - val\_loss: 0.1039 - val\_accuracy: 0.9801  
Epoch 13/20  
469/469 [=====] - 7s 14ms/step - loss: 0.0232 - accuracy: 0.9948 - val\_loss: 0.0865 - val\_accuracy: 0.9825  
Epoch 14/20  
469/469 [=====] - 7s 14ms/step - loss: 0.0251 - accuracy: 0.9941 - val\_loss: 0.1090 - val\_accuracy: 0.9796  
Epoch 15/20  
469/469 [=====] - 7s 14ms/step - loss: 0.0236 - accuracy: 0.9948 - val\_loss: 0.0997 - val\_accuracy: 0.9834  
Epoch 16/20  
469/469 [=====] - 7s 14ms/step - loss: 0.0268 - accuracy: 0.9942 - val\_loss: 0.0856 - val\_accuracy: 0.9832  
Epoch 17/20  
469/469 [=====] - 7s 14ms/step - loss: 0.0295 - accuracy: 0.9938 - val\_loss: 0.0986 - val\_accuracy: 0.9814  
Epoch 18/20  
469/469 [=====] - 6s 14ms/step - loss: 0.0226 - accuracy: 0.9947 - val\_loss: 0.1006 - val\_accuracy: 0.9835  
Epoch 19/20  
469/469 [=====] - 7s 15ms/step - loss: 0.0258 - accuracy: 0.9938 - val\_loss: 0.0886 - val\_accuracy: 0.9835  
Epoch 20/20  
469/469 [=====] - 7s 14ms/step - loss: 0.0265 - accuracy: 0.9940 - val\_loss: 0.1049 - val\_accuracy: 0.9800  
Scores for noise scale 0.1

Test loss: 0.10485726594924927  
Test accuracy: 0.9800000190734863  
Model: "sequential\_5"

Layer (type)	Output Shape	Param #
dense_39 (Dense)	(None, 512)	401920
dropout_26 (Dropout)	(None, 512)	0
dense_40 (Dense)	(None, 512)	262656
dropout_27 (Dropout)	(None, 512)	0
dense_41 (Dense)	(None, 10)	5130
dense_42 (Dense)	(None, 512)	5632
dropout_28 (Dropout)	(None, 512)	0
dense_43 (Dense)	(None, 512)	262656
dropout_29 (Dropout)	(None, 512)	0
dense_44 (Dense)	(None, 10)	5130
dense_45 (Dense)	(None, 512)	5632
dropout_30 (Dropout)	(None, 512)	0
dense_46 (Dense)	(None, 512)	262656
dropout_31 (Dropout)	(None, 512)	0
dense_47 (Dense)	(None, 10)	5130

=====  
Total params: 1,216,542  
Trainable params: 1,216,542  
Non-trainable params: 0

Epoch 1/20  
469/469 [=====] - 17s 28ms/step - loss: 0.0987 - accuracy: 0.9905 - val\_loss: 0.1081 - val\_accuracy: 0.9828  
Epoch 2/20  
469/469 [=====] - 10s 22ms/step - loss: 0.0375 - accuracy: 0.9931 - val\_loss: 0.1096 - val\_accuracy: 0.9828  
Epoch 3/20  
469/469 [=====] - 10s 22ms/step - loss: 0.0337 - accuracy: 0.9934 - val\_loss: 0.1138 - val\_accuracy: 0.9810  
Epoch 4/20  
469/469 [=====] - 10s 21ms/step - loss: 0.0370 - accuracy: 0.9926 - val\_loss: 0.1123 - val\_accuracy: 0.9810  
Epoch 5/20  
469/469 [=====] - 10s 22ms/step - loss: 0.0380 - accuracy: 0.9932 - val\_loss: 0.1454 - val\_accuracy: 0.9790  
Epoch 6/20  
469/469 [=====] - 10s 22ms/step - loss: 0.0396 - accuracy: 0.9928 - val\_loss: 0.1222 - val\_accuracy: 0.9808  
Epoch 7/20  
469/469 [=====] - 10s 21ms/step - loss: 0.0434 - accuracy: 0.9923 - val\_loss: 0.1203 - val\_accuracy: 0.9810  
Epoch 8/20  
469/469 [=====] - 10s 22ms/step - loss: 0.0484 - accuracy: 0.9912 - val\_loss: 0.1192 - val\_accuracy: 0.9807  
Epoch 9/20  
469/469 [=====] - 10s 22ms/step - loss: 0.0406 - accuracy: 0.9928 - val\_loss: 0.1234 - val\_accuracy: 0.9814  
Epoch 10/20  
469/469 [=====] - 10s 22ms/step - loss: 0.0434 - accuracy: 0.9926 - val\_loss: 0.1183 - val\_accuracy: 0.9807  
Epoch 11/20  
469/469 [=====] - 11s 23ms/step - loss: 0.0494 - accuracy: 0.9918 - val\_loss: 0.1174 - val\_accuracy: 0.9812  
Epoch 12/20  
469/469 [=====] - 10s 21ms/step - loss: 0.0492 - accuracy: 0.9915 - val\_loss: 0.1362 - val\_accuracy: 0.9781  
Epoch 13/20  
469/469 [=====] - 10s 22ms/step - loss: 0.0458 - accuracy: 0.9916 - val\_loss: 0.1240 - val\_accuracy: 0.9811  
Epoch 14/20  
469/469 [=====] - 10s 21ms/step - loss: 0.0531 - accuracy: 0.9901 - val\_loss: 0.1192 - val\_accuracy: 0.9791  
Epoch 15/20  
469/469 [=====] - 10s 21ms/step - loss: 0.0525 - accuracy: 0.9898 - val\_loss: 0.1186 - val\_accuracy: 0.9817

Epoch 16/20  
 469/469 [=====] - 10s 21ms/step - loss: 0.0459 - accuracy: 0.9922 - val\_loss: 0.1193 - val\_accuracy: 0.9807  
 Epoch 17/20  
 469/469 [=====] - 10s 20ms/step - loss: 0.0457 - accuracy: 0.9924 - val\_loss: 0.1032 - val\_accuracy: 0.9832  
 Epoch 18/20  
 469/469 [=====] - 9s 20ms/step - loss: 0.0502 - accuracy: 0.9913 - val\_loss: 0.1132 - val\_accuracy: 0.9823  
 Epoch 19/20  
 469/469 [=====] - 10s 21ms/step - loss: 0.0452 - accuracy: 0.9921 - val\_loss: 0.1193 - val\_accuracy: 0.9816  
 Epoch 20/20  
 469/469 [=====] - 10s 21ms/step - loss: 0.0495 - accuracy: 0.9915 - val\_loss: 0.1204 - val\_accuracy: 0.9807  
 Scores for noise scale 0.5  
 Test loss: 0.12037938088178635  
 Test accuracy: 0.9807000160217285  
 Model: "sequential\_5"

Layer (type)	Output Shape	Param #
dense_39 (Dense)	(None, 512)	401920
dropout_26 (Dropout)	(None, 512)	0
dense_40 (Dense)	(None, 512)	262656
dropout_27 (Dropout)	(None, 512)	0
dense_41 (Dense)	(None, 10)	5130
dense_42 (Dense)	(None, 512)	5632
dropout_28 (Dropout)	(None, 512)	0
dense_43 (Dense)	(None, 512)	262656
dropout_29 (Dropout)	(None, 512)	0
dense_44 (Dense)	(None, 10)	5130
dense_45 (Dense)	(None, 512)	5632
dropout_30 (Dropout)	(None, 512)	0
dense_46 (Dense)	(None, 512)	262656
dropout_31 (Dropout)	(None, 512)	0
dense_47 (Dense)	(None, 10)	5130
dense_48 (Dense)	(None, 512)	5632
dropout_32 (Dropout)	(None, 512)	0
dense_49 (Dense)	(None, 512)	262656
dropout_33 (Dropout)	(None, 512)	0
dense_50 (Dense)	(None, 10)	5130

=====  
 Total params: 1,489,960  
 Trainable params: 1,489,960  
 Non-trainable params: 0

Epoch 1/20  
 469/469 [=====] - 16s 26ms/step - loss: 0.1098 - accuracy: 0.9892 - val\_loss: 0.1212 - val\_accuracy: 0.9812  
 Epoch 2/20  
 469/469 [=====] - 12s 25ms/step - loss: 0.0532 - accuracy: 0.9913 - val\_loss: 0.1149 - val\_accuracy: 0.9822  
 Epoch 3/20  
 469/469 [=====] - 12s 26ms/step - loss: 0.0543 - accuracy: 0.9911 - val\_loss: 0.1098 - val\_accuracy: 0.9831  
 Epoch 4/20  
 469/469 [=====] - 12s 26ms/step - loss: 0.0488 - accuracy: 0.9923 - val\_loss: 0.1413 - val\_accuracy: 0.9796  
 Epoch 5/20  
 469/469 [=====] - 12s 25ms/step - loss: 0.0693 - accuracy: 0.9889 - val\_loss: 0.1355 - val\_accuracy: 0.9777  
 Epoch 6/20  
 469/469 [=====] - 12s 25ms/step - loss: 0.0715 - accuracy: 0.9890 - val\_loss: 0.1320 - val\_accuracy: 0.9805  
 Epoch 7/20

```

469/469 [=====] - 11s 24ms/step - loss: 0.0702 - accuracy: 0.9891 - val_loss: 0.1543 - val_accuracy: 0.97
67
Epoch 8/20
469/469 [=====] - 11s 24ms/step - loss: 0.1070 - accuracy: 0.9821 - val_loss: 0.1369 - val_accuracy: 0.97
99
Epoch 9/20
469/469 [=====] - 11s 24ms/step - loss: 0.0909 - accuracy: 0.9852 - val_loss: 0.1346 - val_accuracy: 0.97
76
Epoch 10/20
469/469 [=====] - 12s 26ms/step - loss: 0.0952 - accuracy: 0.9844 - val_loss: 0.1311 - val_accuracy: 0.97
92
Epoch 11/20
469/469 [=====] - 12s 26ms/step - loss: 0.0980 - accuracy: 0.9847 - val_loss: 0.1539 - val_accuracy: 0.97
52
Epoch 12/20
469/469 [=====] - 12s 26ms/step - loss: 0.1071 - accuracy: 0.9825 - val_loss: 0.1494 - val_accuracy: 0.97
62
Epoch 13/20
469/469 [=====] - 12s 25ms/step - loss: 0.1199 - accuracy: 0.9803 - val_loss: 0.1439 - val_accuracy: 0.97
74
Epoch 14/20
469/469 [=====] - 12s 26ms/step - loss: 0.1339 - accuracy: 0.9776 - val_loss: 0.1786 - val_accuracy: 0.97
13
Epoch 15/20
469/469 [=====] - 12s 27ms/step - loss: 0.1361 - accuracy: 0.9765 - val_loss: 0.1597 - val_accuracy: 0.97
44
Epoch 16/20
469/469 [=====] - 713s 2s/step - loss: 0.1450 - accuracy: 0.9737 - val_loss: 0.1693 - val_accuracy: 0.973
0
Epoch 17/20
469/469 [=====] - 11s 24ms/step - loss: 0.1506 - accuracy: 0.9744 - val_loss: 0.1835 - val_accuracy: 0.97
08
Epoch 18/20
469/469 [=====] - 13s 28ms/step - loss: 0.1241 - accuracy: 0.9797 - val_loss: 0.1605 - val_accuracy: 0.97
41
Epoch 19/20
469/469 [=====] - 13s 27ms/step - loss: 0.1204 - accuracy: 0.9804 - val_loss: 0.1598 - val_accuracy: 0.97
38
Epoch 20/20
469/469 [=====] - 13s 28ms/step - loss: 0.1289 - accuracy: 0.9784 - val_loss: 0.1763 - val_accuracy: 0.97
09
Scores for noise scale 1.0
Test loss: 0.1762864738702774
Test accuracy: 0.9708999991416931
Model: "sequential_5"

```

Layer (type)	Output Shape	Param #
=====		
dense_39 (Dense)	(None, 512)	401920
dropout_26 (Dropout)	(None, 512)	0
dense_40 (Dense)	(None, 512)	262656
dropout_27 (Dropout)	(None, 512)	0
dense_41 (Dense)	(None, 10)	5130
dense_42 (Dense)	(None, 512)	5632
dropout_28 (Dropout)	(None, 512)	0
dense_43 (Dense)	(None, 512)	262656
dropout_29 (Dropout)	(None, 512)	0
dense_44 (Dense)	(None, 10)	5130
dense_45 (Dense)	(None, 512)	5632
dropout_30 (Dropout)	(None, 512)	0
dense_46 (Dense)	(None, 512)	262656
dropout_31 (Dropout)	(None, 512)	0
dense_47 (Dense)	(None, 10)	5130
dense_48 (Dense)	(None, 512)	5632
dropout_32 (Dropout)	(None, 512)	0
dense_49 (Dense)	(None, 512)	262656
dropout_33 (Dropout)	(None, 512)	0



dense_50 (Dense)	(None, 10)	5130
dense_51 (Dense)	(None, 512)	5632
dropout_34 (Dropout)	(None, 512)	0
dense_52 (Dense)	(None, 512)	262656
dropout_35 (Dropout)	(None, 512)	0
dense_53 (Dense)	(None, 10)	5130

=====

Total params: 1,763,378  
Trainable params: 1,763,378  
Non-trainable params: 0

---

Epoch 1/20  
469/469 [=====] - 23s 38ms/step - loss: 0.1960 - accuracy: 0.9743 - val\_loss: 0.1736 - val\_accuracy: 0.9728  
Epoch 2/20  
469/469 [=====] - 15s 31ms/step - loss: 0.1804 - accuracy: 0.9684 - val\_loss: 0.2119 - val\_accuracy: 0.9626  
Epoch 3/20  
469/469 [=====] - 14s 30ms/step - loss: 0.1995 - accuracy: 0.9646 - val\_loss: 0.2096 - val\_accuracy: 0.9601  
Epoch 4/20  
469/469 [=====] - 15s 32ms/step - loss: 0.1991 - accuracy: 0.9649 - val\_loss: 0.1966 - val\_accuracy: 0.9653  
Epoch 5/20  
469/469 [=====] - 15s 33ms/step - loss: 0.1858 - accuracy: 0.9667 - val\_loss: 0.1930 - val\_accuracy: 0.9680  
Epoch 6/20  
469/469 [=====] - 15s 32ms/step - loss: 0.1784 - accuracy: 0.9692 - val\_loss: 0.2196 - val\_accuracy: 0.9613  
Epoch 7/20  
469/469 [=====] - 14s 30ms/step - loss: 0.2086 - accuracy: 0.9629 - val\_loss: 0.1934 - val\_accuracy: 0.9672  
Epoch 8/20  
469/469 [=====] - 14s 30ms/step - loss: 0.1738 - accuracy: 0.9697 - val\_loss: 0.1838 - val\_accuracy: 0.9715  
Epoch 9/20  
469/469 [=====] - 14s 31ms/step - loss: 0.1601 - accuracy: 0.9730 - val\_loss: 0.1747 - val\_accuracy: 0.9715  
Epoch 10/20  
469/469 [=====] - 14s 30ms/step - loss: 0.1677 - accuracy: 0.9717 - val\_loss: 0.2170 - val\_accuracy: 0.9642  
Epoch 11/20  
469/469 [=====] - 14s 30ms/step - loss: 0.2282 - accuracy: 0.9553 - val\_loss: 0.2368 - val\_accuracy: 0.9552  
Epoch 12/20  
469/469 [=====] - 14s 29ms/step - loss: 0.2080 - accuracy: 0.9611 - val\_loss: 0.2294 - val\_accuracy: 0.9561  
Epoch 13/20  
469/469 [=====] - 14s 30ms/step - loss: 0.2651 - accuracy: 0.9469 - val\_loss: 0.2522 - val\_accuracy: 0.9531  
Epoch 14/20  
469/469 [=====] - 14s 30ms/step - loss: 0.2241 - accuracy: 0.9592 - val\_loss: 0.2108 - val\_accuracy: 0.9649  
Epoch 15/20  
469/469 [=====] - 14s 31ms/step - loss: 0.2242 - accuracy: 0.9597 - val\_loss: 0.2303 - val\_accuracy: 0.9581  
Epoch 16/20  
469/469 [=====] - 14s 30ms/step - loss: 0.2224 - accuracy: 0.9597 - val\_loss: 0.1854 - val\_accuracy: 0.9695  
Epoch 17/20  
469/469 [=====] - 14s 30ms/step - loss: 0.2467 - accuracy: 0.9473 - val\_loss: 0.3080 - val\_accuracy: 0.9206  
Epoch 18/20  
469/469 [=====] - 14s 30ms/step - loss: 0.2531 - accuracy: 0.9502 - val\_loss: 0.2386 - val\_accuracy: 0.9601  
Epoch 19/20  
469/469 [=====] - 14s 31ms/step - loss: 0.2523 - accuracy: 0.9517 - val\_loss: 0.2798 - val\_accuracy: 0.9490  
Epoch 20/20  
469/469 [=====] - 14s 29ms/step - loss: 0.2491 - accuracy: 0.9541 - val\_loss: 0.2333 - val\_accuracy: 0.9591  
Scores for noise scale 2.0  
Test loss: 0.233308807015419  
Test accuracy: 0.9591000080108643  
Model: "sequential\_5"

---

Layer (type)	Output Shape	Param #
dense_39 (Dense)	(None, 512)	401920

---

dropout_26 (Dropout)	(None, 512)	0
dense_40 (Dense)	(None, 512)	262656
dropout_27 (Dropout)	(None, 512)	0
dense_41 (Dense)	(None, 10)	5130
dense_42 (Dense)	(None, 512)	5632
dropout_28 (Dropout)	(None, 512)	0
dense_43 (Dense)	(None, 512)	262656
dropout_29 (Dropout)	(None, 512)	0
dense_44 (Dense)	(None, 10)	5130
dense_45 (Dense)	(None, 512)	5632
dropout_30 (Dropout)	(None, 512)	0
dense_46 (Dense)	(None, 512)	262656
dropout_31 (Dropout)	(None, 512)	0
dense_47 (Dense)	(None, 10)	5130
dense_48 (Dense)	(None, 512)	5632
dropout_32 (Dropout)	(None, 512)	0
dense_49 (Dense)	(None, 512)	262656
dropout_33 (Dropout)	(None, 512)	0
dense_50 (Dense)	(None, 10)	5130
dense_51 (Dense)	(None, 512)	5632
dropout_34 (Dropout)	(None, 512)	0
dense_52 (Dense)	(None, 512)	262656
dropout_35 (Dropout)	(None, 512)	0
dense_53 (Dense)	(None, 10)	5130
dense_54 (Dense)	(None, 512)	5632
dropout_36 (Dropout)	(None, 512)	0
dense_55 (Dense)	(None, 512)	262656
dropout_37 (Dropout)	(None, 512)	0
dense_56 (Dense)	(None, 10)	5130

=====

Total params: 2,036,796  
Trainable params: 2,036,796  
Non-trainable params: 0

---

Epoch 1/20  
469/469 [=====] - 22s 38ms/step - loss: 0.3320 - accuracy: 0.9433 - val\_loss: 0.3415 - val\_accuracy: 0.9262

Epoch 2/20  
469/469 [=====] - 17s 37ms/step - loss: 0.4378 - accuracy: 0.8789 - val\_loss: 0.4468 - val\_accuracy: 0.8710

Epoch 3/20  
469/469 [=====] - 15s 31ms/step - loss: 0.3376 - accuracy: 0.9237 - val\_loss: 0.3357 - val\_accuracy: 0.9188

Epoch 4/20  
469/469 [=====] - 16s 34ms/step - loss: 0.3545 - accuracy: 0.9096 - val\_loss: 0.3504 - val\_accuracy: 0.9123

Epoch 5/20  
469/469 [=====] - 16s 34ms/step - loss: 0.5081 - accuracy: 0.8572 - val\_loss: 0.5264 - val\_accuracy: 0.8455

Epoch 6/20  
469/469 [=====] - 16s 34ms/step - loss: 0.5417 - accuracy: 0.8374 - val\_loss: 0.6627 - val\_accuracy: 0.7739

Epoch 7/20  
469/469 [=====] - 16s 35ms/step - loss: 0.6509 - accuracy: 0.8030 - val\_loss: 0.6186 - val\_accuracy: 0.8092

Epoch 8/20  
469/469 [=====] - 15s 33ms/step - loss: 0.6836 - accuracy: 0.7771 - val\_loss: 0.5779 - val\_accuracy: 0.80

```

34
Epoch 9/20
469/469 [=====] - 16s 34ms/step - loss: 0.7000 - accuracy: 0.7410 - val_loss: 0.7051 - val_accuracy: 0.71
66
Epoch 10/20
469/469 [=====] - 15s 31ms/step - loss: 0.7518 - accuracy: 0.6990 - val_loss: 0.8444 - val_accuracy: 0.67
02
Epoch 11/20
469/469 [=====] - 16s 33ms/step - loss: 1.0483 - accuracy: 0.5979 - val_loss: 1.0719 - val_accuracy: 0.58
88
Epoch 12/20
469/469 [=====] - 16s 33ms/step - loss: 1.0148 - accuracy: 0.6141 - val_loss: 0.8322 - val_accuracy: 0.67
91
Epoch 13/20
469/469 [=====] - 16s 34ms/step - loss: 1.2156 - accuracy: 0.5461 - val_loss: 1.3663 - val_accuracy: 0.46
49
Epoch 14/20
469/469 [=====] - 16s 34ms/step - loss: 1.2675 - accuracy: 0.4964 - val_loss: 1.2321 - val_accuracy: 0.48
64
Epoch 15/20
469/469 [=====] - 17s 35ms/step - loss: 1.1191 - accuracy: 0.5296 - val_loss: 1.0299 - val_accuracy: 0.57
07
Epoch 16/20
469/469 [=====] - 16s 34ms/step - loss: 1.0752 - accuracy: 0.5448 - val_loss: 1.0453 - val_accuracy: 0.57
15
Epoch 17/20
469/469 [=====] - 15s 32ms/step - loss: 0.9758 - accuracy: 0.6197 - val_loss: 0.9282 - val_accuracy: 0.64
66
Epoch 18/20
469/469 [=====] - 15s 32ms/step - loss: 1.1279 - accuracy: 0.5638 - val_loss: 1.2162 - val_accuracy: 0.52
56
Epoch 19/20
469/469 [=====] - 15s 33ms/step - loss: 1.2785 - accuracy: 0.4974 - val_loss: 1.1748 - val_accuracy: 0.53
35
Epoch 20/20
469/469 [=====] - 16s 34ms/step - loss: 1.2989 - accuracy: 0.5041 - val_loss: 1.2653 - val_accuracy: 0.50
87
Scores for noise scale 4.0
Test loss: 1.2652525901794434
Test accuracy: 0.5087000131607056

```

As expected increased noise increases losses and decreases accuracy

```


In [110... losses = [0.0777, 0.1049, 0.1204, 0.1763, 0.2333, 1.2653]
accuracy = [0.9839, 0.9800, 0.9807, 0.9709, 0.9591, 0.5087]

```

```

In [113... plt.figure()
plot = plt.plot(scale, losses)
plt.ylim([0, 1.4])
plt.xlabel('Noisiness')
plt.ylabel('Losses')
plt.show()

```




```

In [115... plt.figure()
plot = plt.plot(scale, accuracy)
plt.ylim([0, 1])

```

```
plt.xlabel('Noisiness')  
plt.ylabel('Accuracy')  
plt.show()
```



In [ ]:

